



CloudButton

Serverless FaaS

Pedro Garcia Lopez

Table of Contents

- CloudLab Research group
- A bit of history
- What is serverless?
- FaaS Orchestration
- FaaS Architecture
- Serverless Data Analytics

CLOUDLAB

Who am I

- Associate Professor at Universitat Rovira i Virgili
- Head of the Cloud and Distributed Systems Lab (2003-2019)
- My research interests and experience: Distributed Systems, Middleware, Storage systems, Cloud Computing, Data Analytics
- Principal Investigator of large research projects (>10)
- 20 years teaching distributed systems and Advanced Programming Techniques

CLOUDLAB

CloudLab: Cloud and Distributed Systems Research group

- **Professors:** Pedro Garcia Lopez, Marc Sanchez Artigas, Carlos Molina Clemente, Carlos Aliagas Castells
- **Postdocs:** Josep Sampé Domenech
- **Senior Research Engineer:** Gerard Paris
- **PhDs:** Daniel Barcelona, Ammar Okran, Alvaro Ruiz
- **Master Students:** Amanda Gómez
- **Projects:**
 - H2020 CloudButton,
 - Software Defined Edge Clouds

CLOUDLAB

Serverless Data Analytics



CloudButton

- 4.4M€ European research project
- cloudbutton.eu
- Coordinated by URV
- 2019-2021





Imperial College
London



AtoS



Why Serverless



WHY SERVERLESS


Blockchains or Cloud ?

- When P2P declined this was an important question
- BitCoin: From 0 to 20K\$ to 3K\$
- Ethereum, ICOs, Smart Contracts
- Bram Cohen sold Bittorrent.com to Tron for 140M\$



primeras compras de criptomonedas

198 visualizaciones •
Hace 3 meses



QUE ES Y COMO FUNCIONA BLOCKCHAIN

320 visualizaciones •
Hace 3 meses



AL PRINCIPIO ERA LA NADA

384 visualizaciones •
Hace 4 meses



SEGUNDA PRESENTACION


637 visualizaciones •
Hace 4 meses



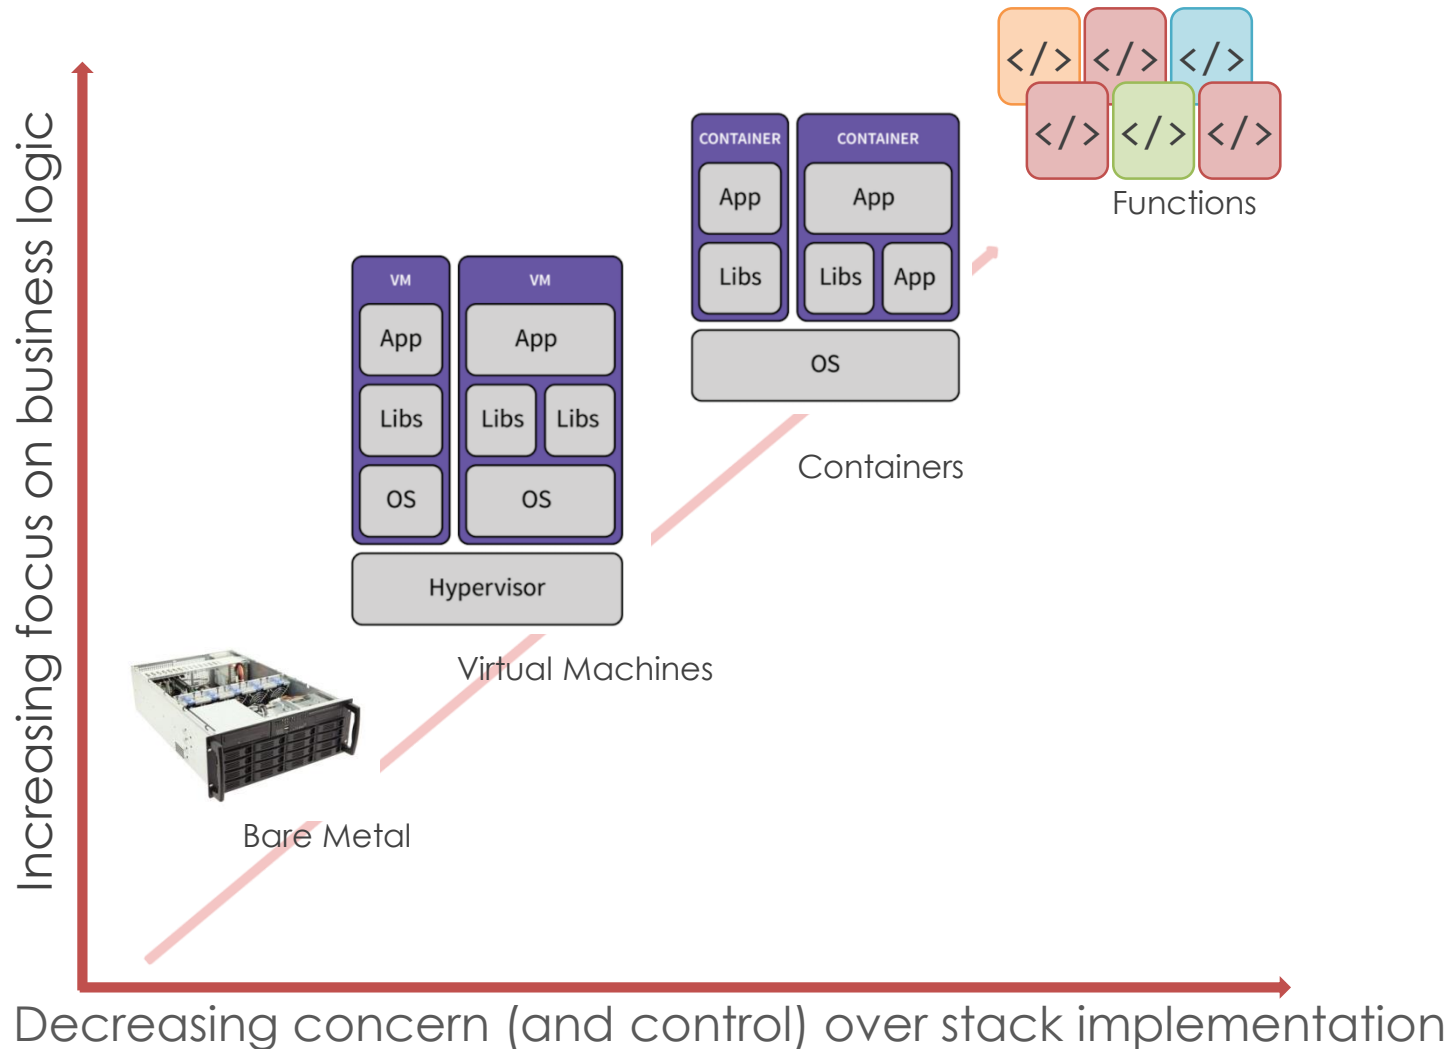
MUNDO FINANCIERO Y CRIPTOMONEDAS

265 visualizaciones •
Hace 4 meses

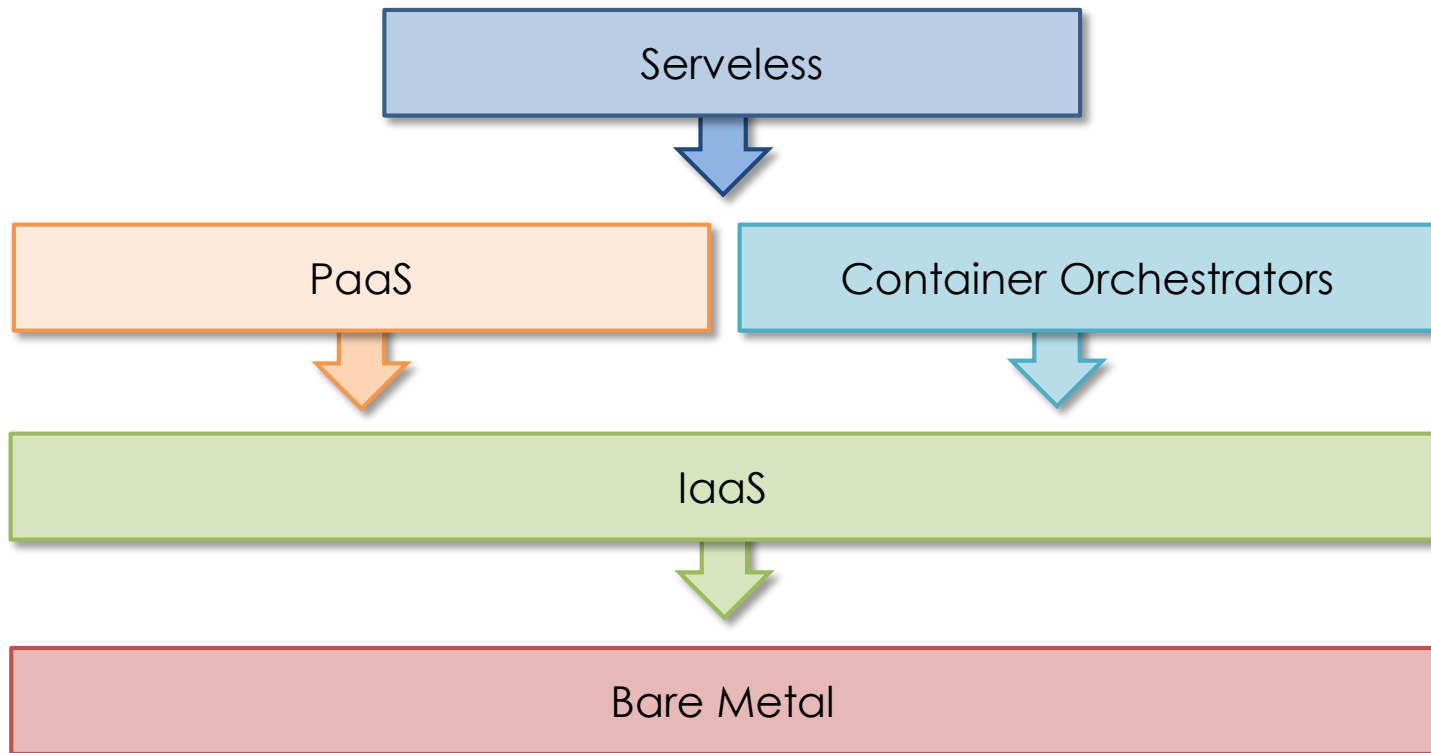
The Monolith



Cloud Computing Evolution




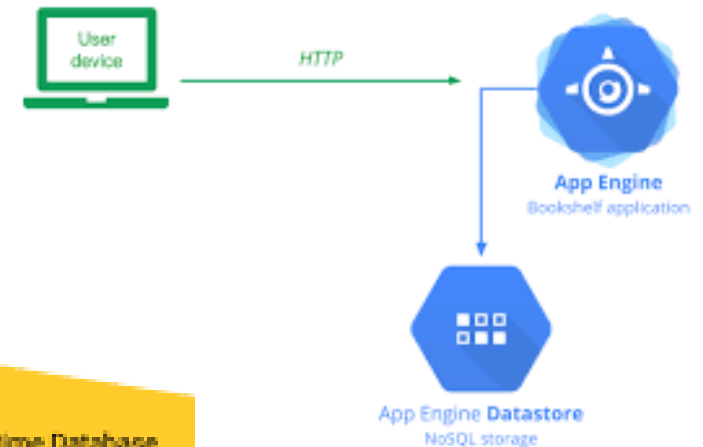
Enter Serverless



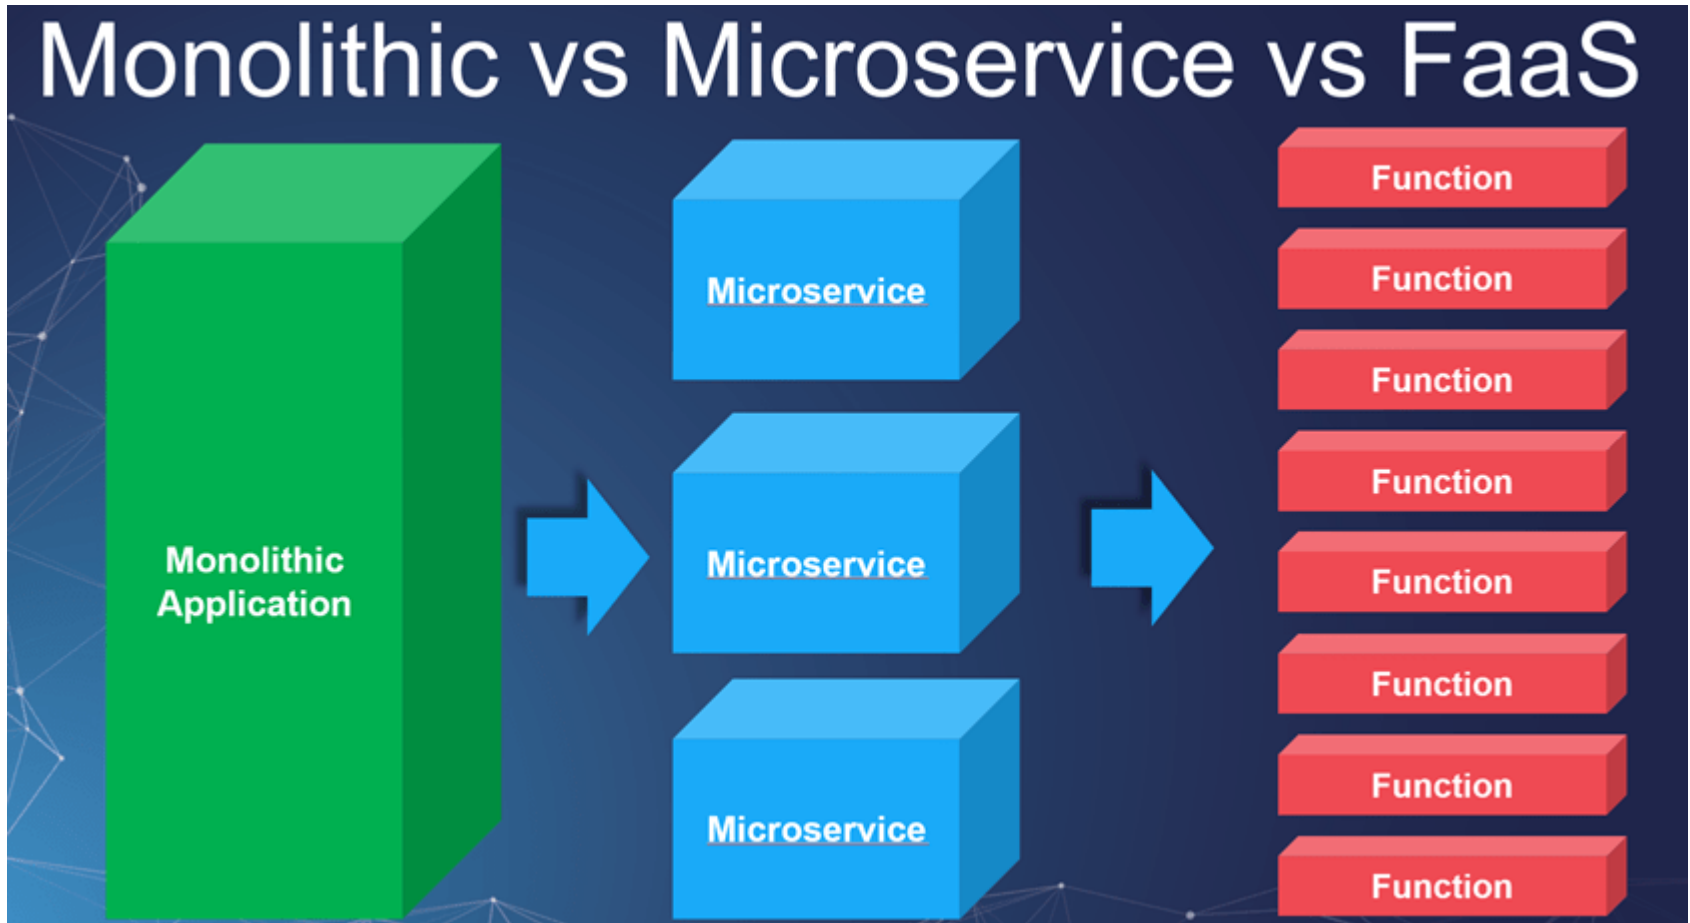
Serverless origins



- Google App Engine PaaS vs Amazon IaaS
- Parse.com
- Backend as a Service
 - Firebase



Cloud Computing Evolution



SERVERLESS COMPUTING

What is Serverless?



As FAAS (Function-as-a-Service):

a cloud-native platform

FOR

short-running, stateless computation

AND

event- and (data-driven) applications

WHICH

scales up and down instantly and automatically

AND

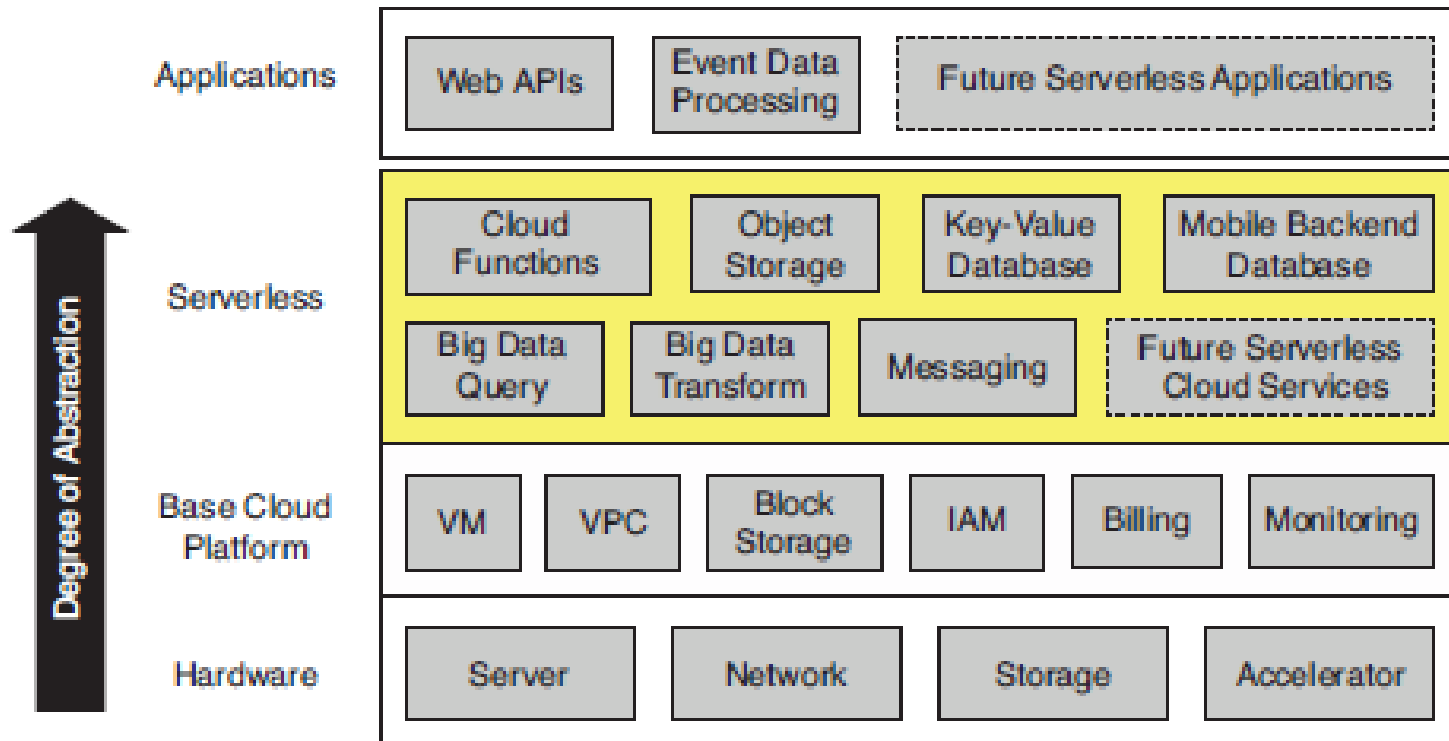
charges for actual usage at a millisecond granularity

What is Serverless?



	<i>Characteristic</i>	<i>AWS Serverless Cloud</i>	<i>AWS Serverful Cloud</i>
PROGRAMMER	When the program is run	On event selected by Cloud user	Continuously until explicitly stopped
	Programming Language	JavaScript, Python, Java, Go, C#, etc. ⁴	Any
	Program State	Kept in storage (stateless)	Anywhere (stateful or stateless)
	Maximum Memory Size	0.125 - 3 GiB (Cloud user selects)	0.5 - 1952 GiB (Cloud user selects)
	Maximum Local Storage	0.5 GiB	0 - 3600 GiB (Cloud user selects)
	Maximum Run Time	900 seconds	None
	Minimum Accounting Unit	0.1 seconds	60 seconds
	Price per Accounting Unit	\$0.00000002 (assuming 0.125 GiB)	\$0.0000867 - \$0.4080000
	Operating System & Libraries	Cloud provider selects ⁵	Cloud user selects
SYSADMIN	Server Instance	Cloud provider selects	Cloud user selects
	Scaling ⁶	Cloud provider responsible	Cloud user responsible
	Deployment	Cloud provider responsible	Cloud user responsible
	Fault Tolerance	Cloud provider responsible	Cloud user responsible
	Monitoring	Cloud provider responsible	Cloud user responsible
	Logging	Cloud provider responsible	Cloud user responsible

What is Serverless?



What is Serverless?

- **Function** (“Action”)
 - Containerized custom-written application code
 - Should include bundled dependencies & binaries
 - Memory & execution time limits
- **Triggers** (“Events”)
 - Causes function execution
 - Can be another function
 - Examples:
 - Upload of a video or image
 - Git commit to a repository
 - ...
- **Resources**
 - External BaaS/PaaS/FaaS services (object storage, queueing, elastic cache, etc.)

What is Serverless?

- **Function** example
(pandas and numpy are dependencies)

```
import pandas as pd
import numpy as np

def main(args):
    dates = pd.date_range('20130101', periods=2)
    df = pd.DataFrame(np.random.randn(2,2), index=dates,
columns=list('AB'))
    print(df)
    return df.to_dict('split')
```

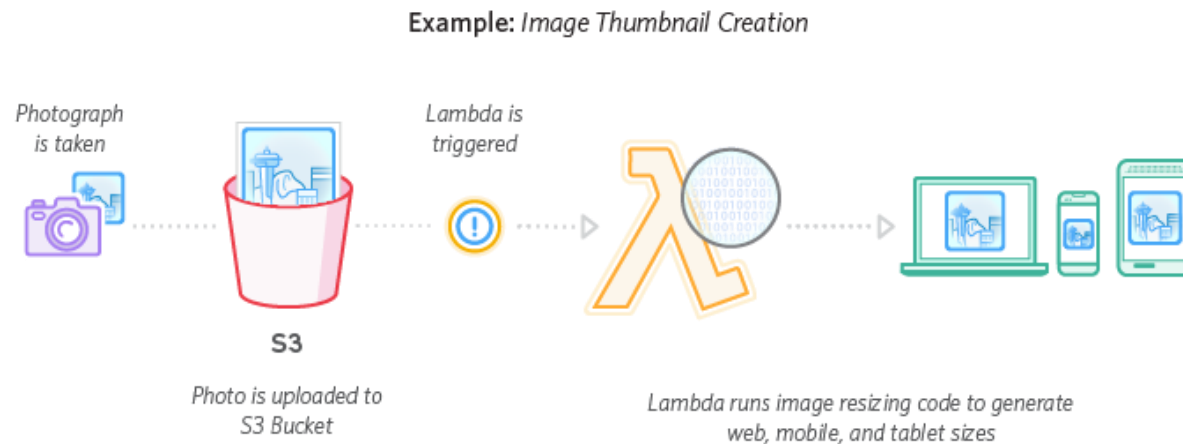
In [12]: df

Out[12]:

	A	B
2013-01-01	0.468173	0.64710
2013-01-02	-0.297858	-0.07476

Serverless Pattern

- An application is architected as a set of business logic **functions**, triggered by discrete **events** or **requests**



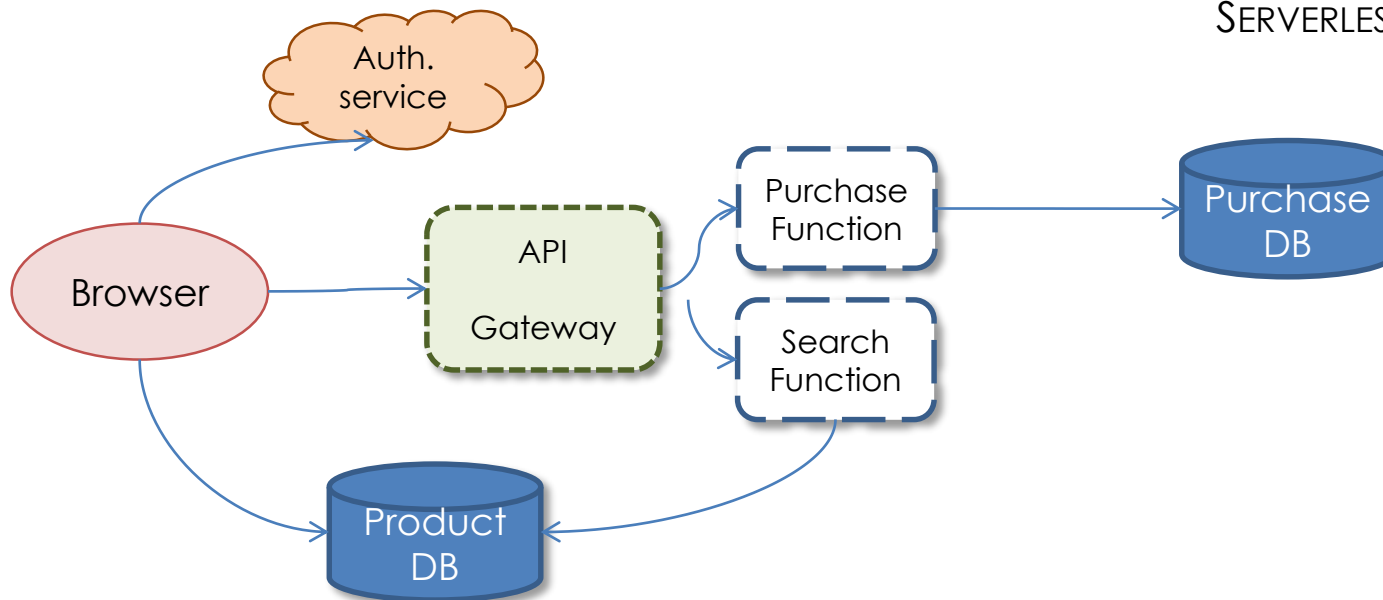
- Good** for microservices, IoT, modest stream processing, ML inferencing, etc.

Serverless Pattern

- Another good example is a typical e-commerce app



TRADITIONAL
SERVERLESS



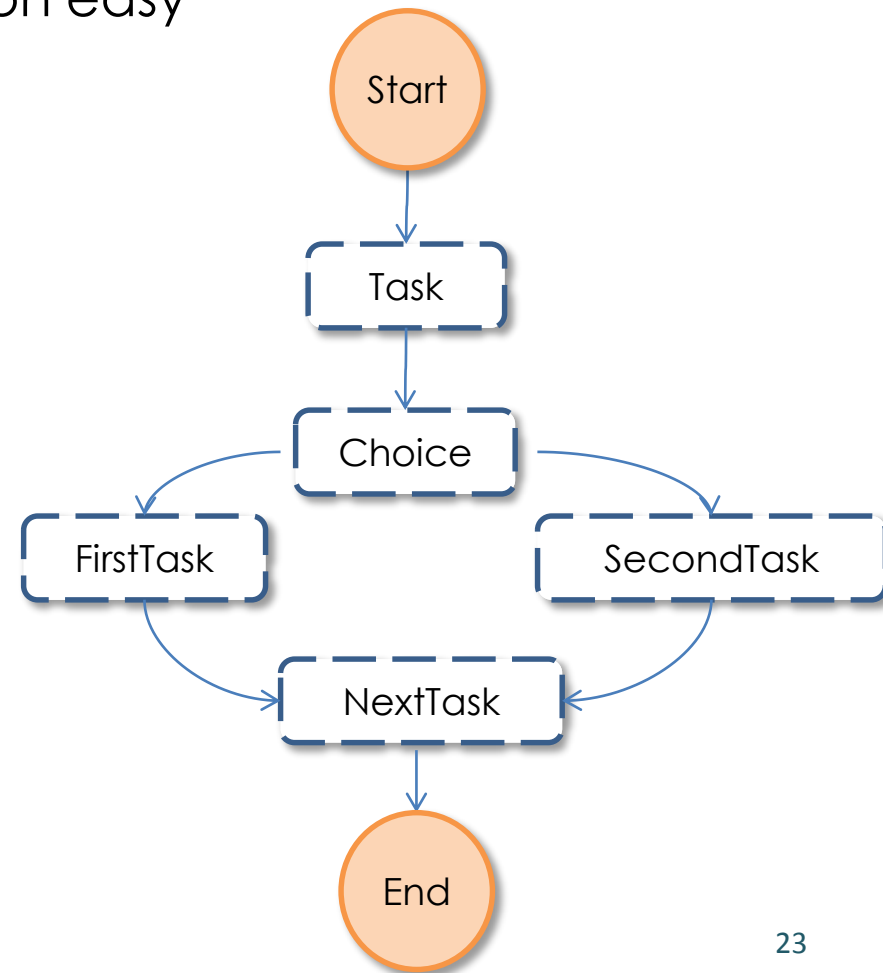
Serverless Pattern



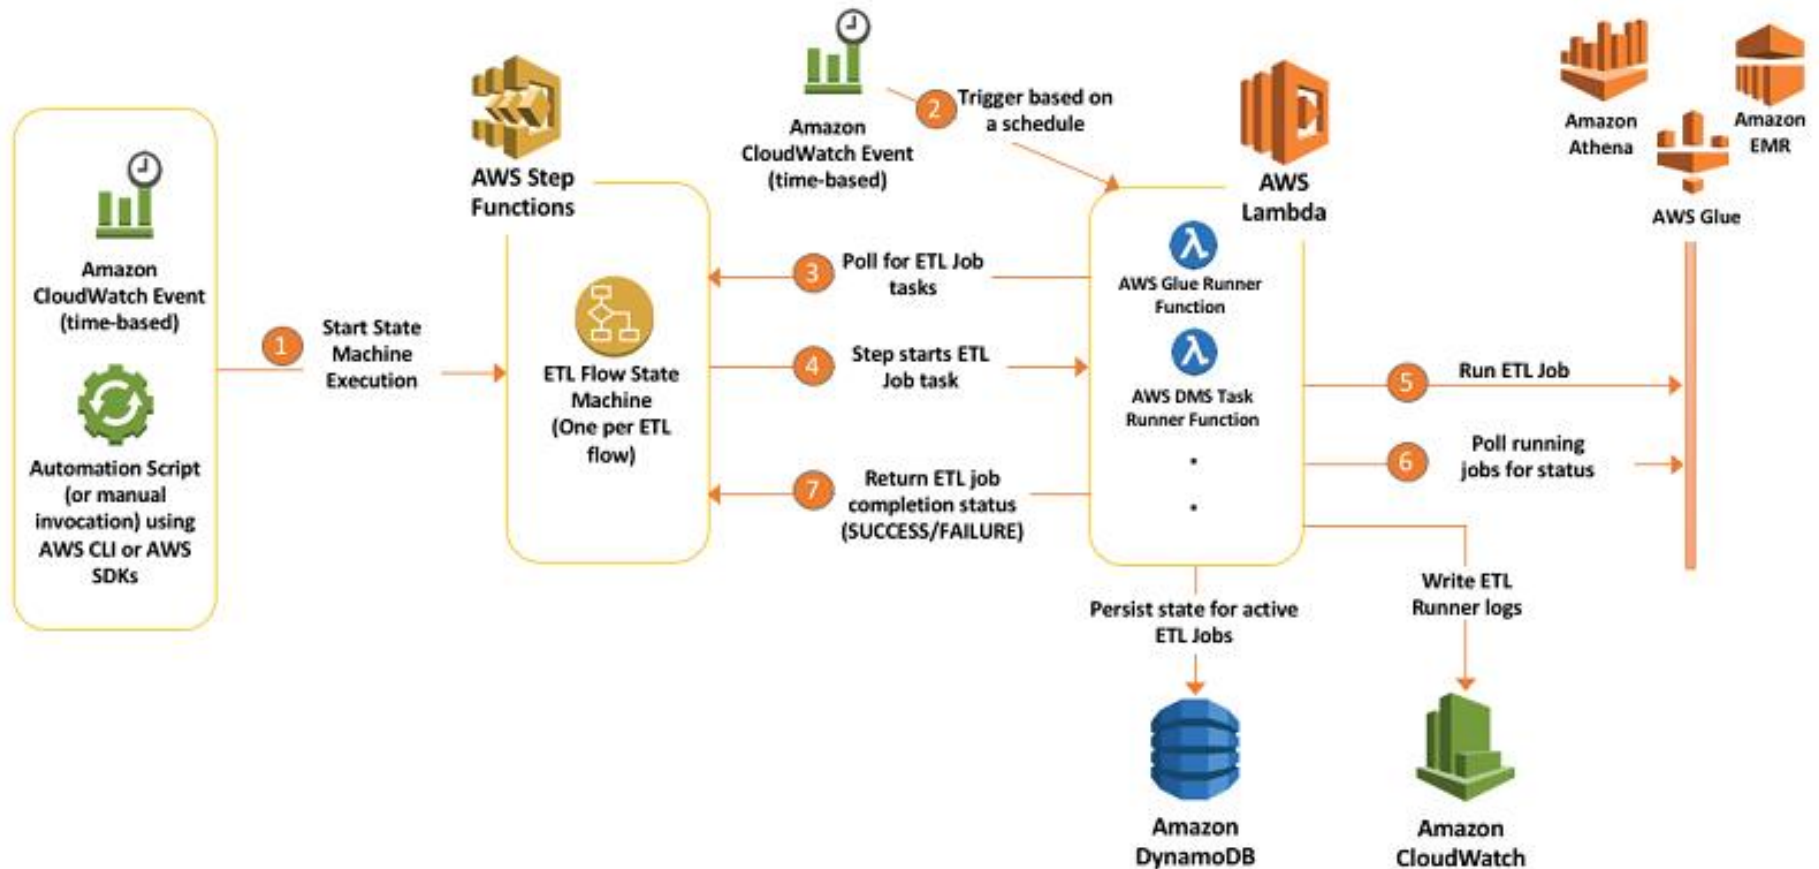
AWS Step Functions

- Also, there is tools such as **AWS Step Functions** that make function and workflow orchestration easy
- State Example: **Choice**

```
“Choice”: {  
  “Type”: “Choice”,  
  “Choices”: [  
    {  
      “Variable”: “$.foo”,  
      “NumericEquals”: 1,  
      “Next”: “FirstTask”  
    },  
    {  
      “Variable”: “$.foo”,  
      “NumericEquals”: 2,  
      “Next”: “SecondTask”  
    }  
  ]  
}
```



Serverless Pattern



Applications



<i>Percent</i>	<i>Use Case</i>
32%	Web and API serving
21%	Data Processing, e.g., batch ETL (database Extract, Transform, and Load)
17%	Integrating 3rd Party Services
16%	Internal tooling
8%	Chat bots e.g., Alexa Skills (SDK for Alexa AI Assistant)
6%	Internet of Things


Table 4: Popularity of serverless computing use cases according to a 2018 survey [22].

Why is Serverless (Un)attractive?

	On-premise	VMs	Containers	Serverless
Time to provision	Weeks-Months	Minutes	Seconds-minutes	Milliseconds
Utilization	Low	High	Higher	Highest
Charging granularity	CapEx	Hours	Minutes	Interval of milliseconds

- The **Good**
 - Removal of the need for a traditional always-on servers
 - Making app development dramatically faster, cheaper, easier
 - Highly available and scalable apps with zero administration
- The **Bad**
 - No in-server state for serverless functions
 - Limited computation times and memory can entail app refactoring
 - Functions are not directly network-addressable

Why is Serverless Unattractive?




Why is Serverless Unattractive?

- FaaS is a Data-Shipping Architecture
- FaaS Stymies Distributed Computing (addressability, communication)
- FaaS stymies hardware-accelerated software innovation
- FaaS discourages Open Source service innovation.
- Limited Lifetimes
- Communication Through Slow Storage
- No guarantees, no QoS
- It can be expensive
- Cold starts !


FAAS ORCHESTRATION

Creating Serverless Workflows




USER FUNCTIONS

WORKFLOW FRAMEWORK



AWS Step Functions



Azure Durable Functions



IBM Function Composer

Amazon Step Functions



AWS Step Functions

```
StateMachine.Builder stateMachineBuilder =
    stateMachine()
        .comment("A_state_machine_with_par_states.")
        .startAt("Parallel");

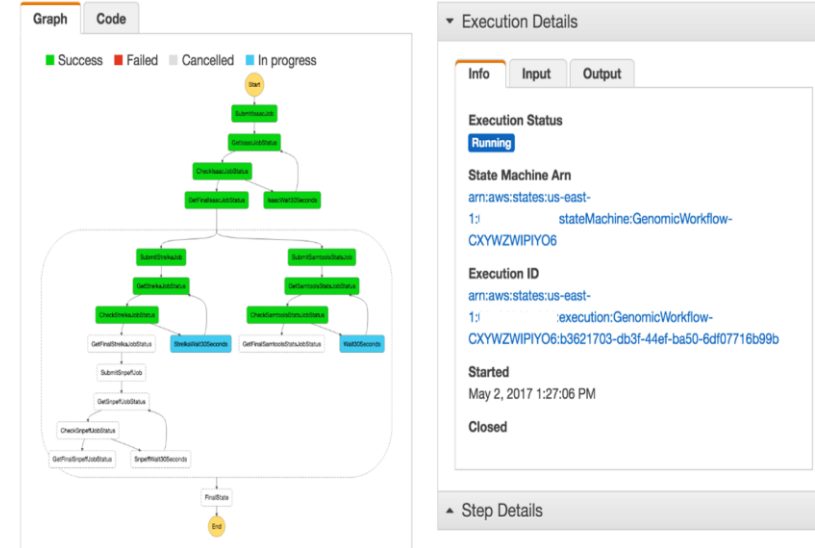
Branch.Builder[] branchBuilders =
    new Branch.Builder[NSTEPS];

for (int i = 0; i < NSTEPS; i++) {
    branchBuilders[i] = branch()
        .startAt(String.valueOf(i + 1))
        .state(String.valueOf(i + 1),
            taskState()
                .resource(arnTask).transition(end()));
}

stateMachineBuilder.state("Parallel",
    parallelState().branches(branchBuilders)
        .transition(end()));
final StateMachine stateMachine =
    stateMachineBuilder.build();
```

b3621703-db3f-44ef-ba50-6df07716b99b

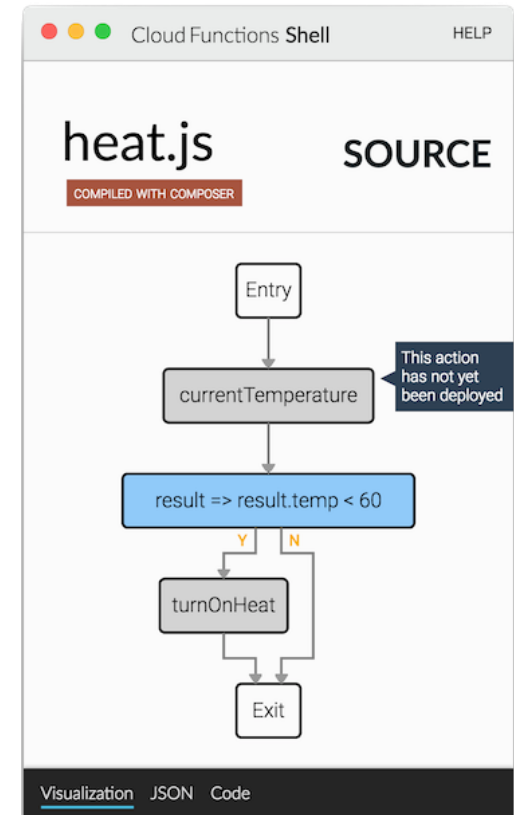
Stop execution




IBM Composer



```
composer.sequence(  
  // programmatic composition  
  'currentTemperature',  
  // call cloud function or API  
  composer.if(  
    // conditional control flow  
    result => result.temp < 60,  
    // mix inline JavaScript  
    'turnOnHeat')  
  // interface to 3rd party services  
)
```



Azure Durable Functions



```
public static async Task Run(DurableOrchestrationContext ctx)
{
    var parallelTasks = new List<Task<int>>();

    // get a list of N work items to process in parallel
    object[] workBatch = await ctx.CallActivityAsync<object[]>("F1");
    for (int i = 0; i < workBatch.Length; i++)
    {
        Task<int> task = ctx.CallActivityAsync<int>("F2", workBatch[i]);
        parallelTasks.Add(task);
    }

    await Task.WhenAll(parallelTasks);

    // aggregate all N outputs and send result to F3
    int sum = parallelTasks.Sum(t => t.Result);
    await ctx.CallActivityAsync("F3", sum);
}
```

Sequences




```
StateMachine.Builder stateMachineBuilder =  
    stateMachine()  
        .comment("A Sequence state machine")  
        .startAt("1");  
for (int i = 1; i <= NSTEPS; i++) {  
    stateMachineBuilder.state(String.valueOf(i),  
        taskState().resource(arnTask)  
        .transition((i != NSTEPS) ?  
            next(String.valueOf(i + 1)) : end()));  
}  
StateMachine stateMachine =  
    stateMachineBuilder.build();
```

```
composer.repeat(40, 'sleepAction')
```


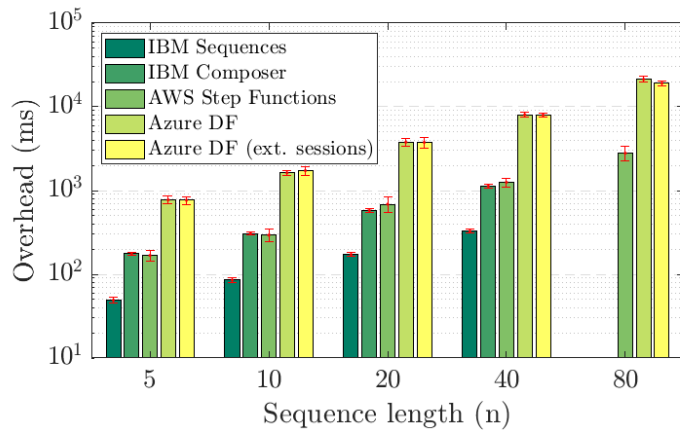
```
for (int i = 0; i < NSTEPS; i++) {  
    await context.  
        CallActivityAsync("sleepAction", null);  
}
```



AWS Step Functions



Sequences

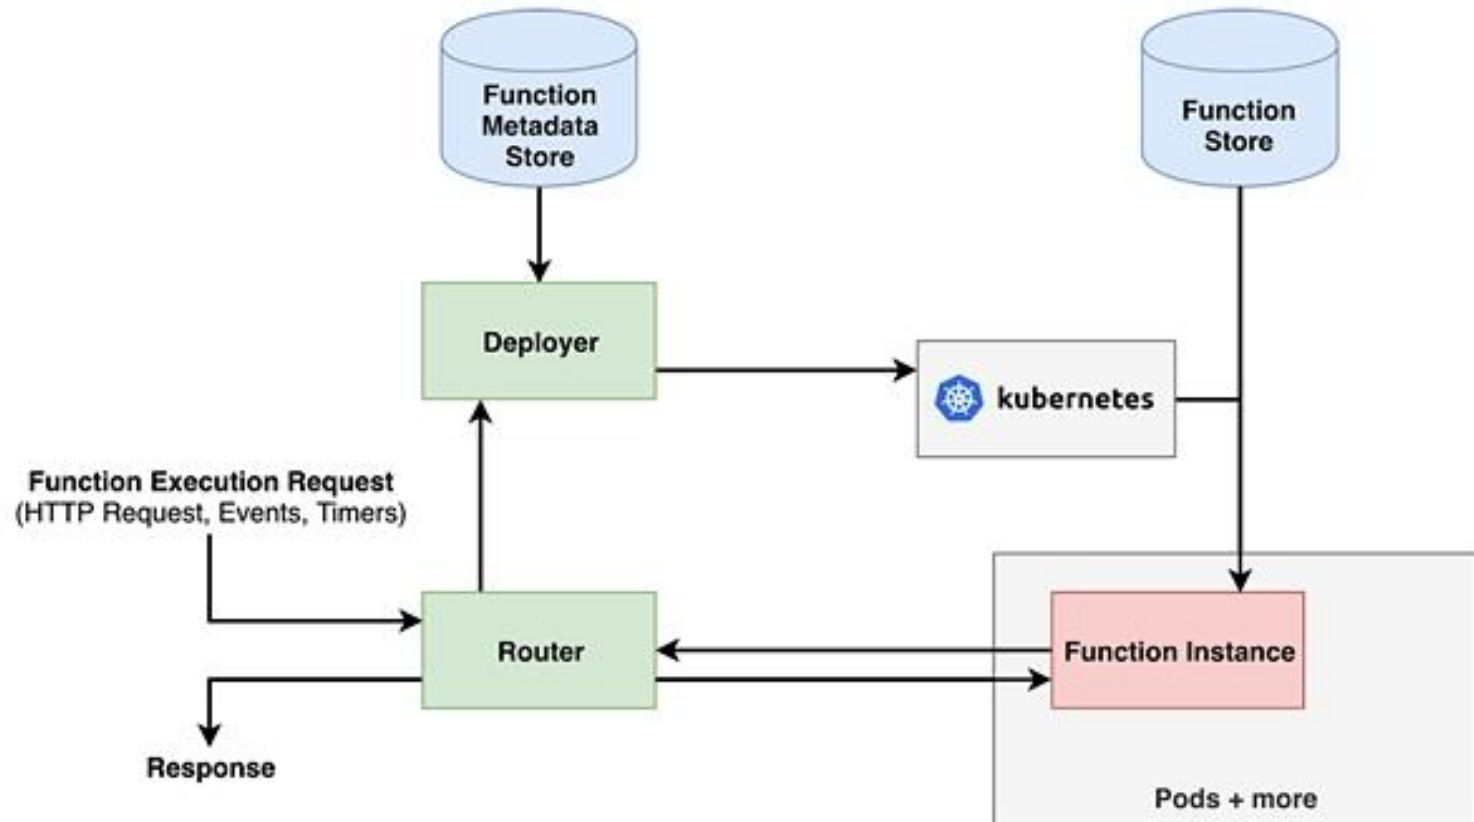


Comparison

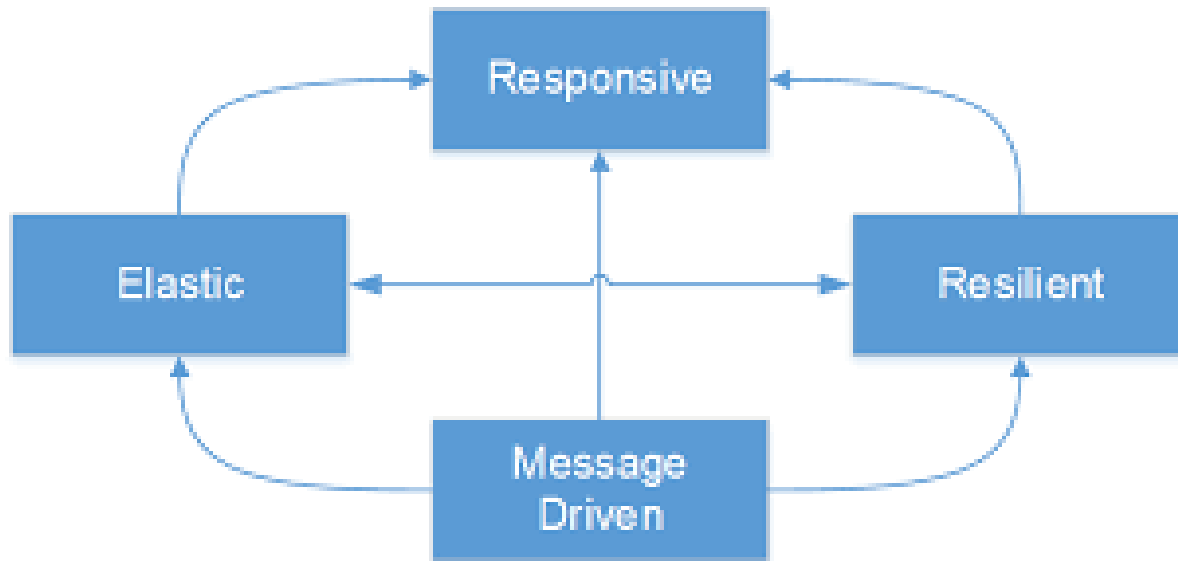
- Amazon Step Functions is the most mature project
- Microsoft ADF is the more advanced in programmability, IBM Composer wins in simplicity
- None of them support parallel tasks efficiently
- Orchestration must have a cost if it is fault-tolerant
- Early immature projects with high potential for the future

FAAS ARCHITECTURE

Architecture



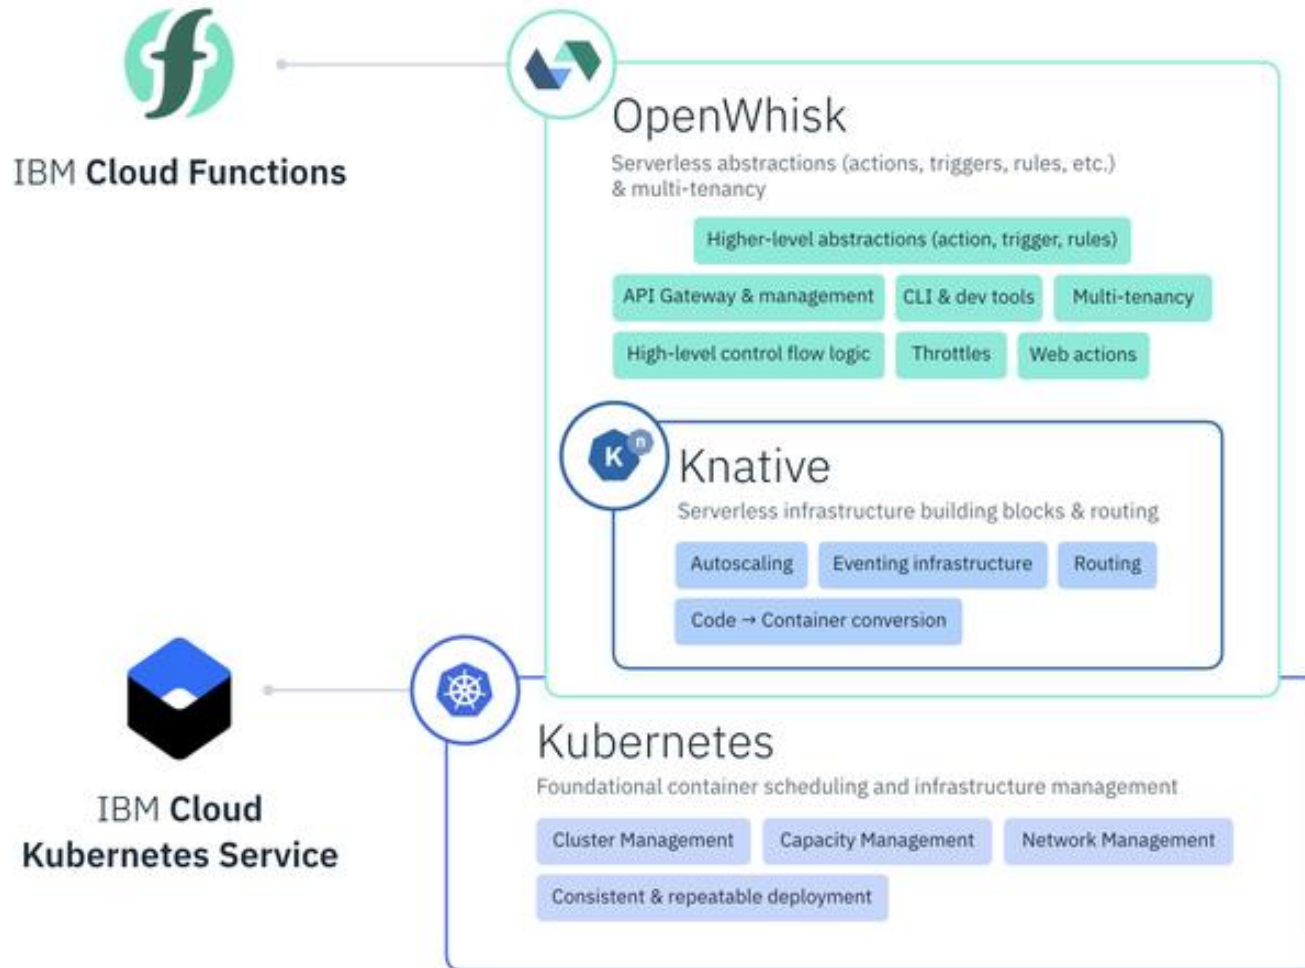
Reactive Core



Knative



CLOUD NATIVE COMPUTING FOUNDATION




Service mesh



CLOUD NATIVE
COMPUTING FOUNDATION


Cloud Native Apps (Microservices) – The New Stack



SERVERLESS DATA ANALYTICS

Serverless Data Analytics?

- Abide by the **functional** programming paradigm:
 - Embarrassingly parallel functions
 - Immutable data through “slow” storage (e.g., S3)
 - **PyWren**[†] and **ExCamera**[‡] research projects show that functions can perform a wider variety of such “**map**” functions
 - **PyWren**[†]’s word count job on **83M** items is only **17%** slower than PySpark running on dedicated servers



[†] Occupy the Cloud: Distributed Computing for the 99%. ACM SOCC 2017

[‡] Encoding, Fast and Slow: Low-Latency Video Processing Using Thousands of Tiny Threads. USENIX NSDI 2017

Serverless Data Analytics?

- One can do a lot of things with a `map(function, data)`

```
def addone(x):  
    return x + 1
```

```
wrenexec= pywren.default_executor()  
data = range(1, 10)  
futures = wrenexec.map(addone, data)
```

Output: [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

- **Functional, declarative** programming models simplify consistency and fault tolerance
- Domain experts tend to write imperative programs
 - Java, Matlab, C++, R, Python, Fortran, ...
- Mismatch between experts' coding skills and analytics

Why CloudButton ?

“Our proposal in this paper was motivated by a professor of computer graphics at UC Berkeley asking us “Why is there no **cloud button**?” He outlined how his students simply wish they could easily “push a button” and have their code – existing, optimized, single-machine code – running on the cloud.”




Figure 1: System architecture for stateless functions

Occupy the Cloud: Distributed Computing for the 99%.
SOCC 2017.

Impact

- Serverless Data Analytics:

- Much simpler for data scientists than analytics frameworks
- Data lakes and FaaS
- Run “single threaded” code in parallel
- Use Cases: Data lakes, Data pipelines, Data flows, IoT



- In-Memory Computing:

- In the industry: Apache Spark ecosystem
- Much better than classical MapReduce for iterative processing
- Operate on datasets in memory to run faster computations
- Use cases: Sorting, clustering, regressions, interactive queries,...



Objectives

- The main goal is to create CloudButton: a Serverless Data Analytics Platform. CloudButton will “democratize big data” by overly simplifying the overall life cycle and programming model thanks to serverless technologies.
- To demonstrate the impact of the project, we target two settings with large data volumes: bioinformatics (genomics, metabolomics) and geospatial data (LiDAR, satellital).



CloudButton



Imperial College
London



Atos



THANK YOU!



This project has received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No 825184.

What is next ?

FaaS and Furious by Forrest Brazeal



"Hello, and welcome to my recently retitled breakout session: 'My Cool Product That Became Irrelevant When AWS Released a New Service Right Before This Breakout Session'."